Allow user overriding of `strong_order` in P0768R1

Gašper Ažman

2018-01-06

Contents

1 Status of this paper

This paper is a defect-report to a paper that has been voted into the working draft. It seeks to highlight an issue with the currently proposed `strong_order` algorithm.

The wording for the entire fix is not provided in this paper, and shall be written if this paper receives support.

2 Problem description

The paper P0768R1[1] proposes the library extensions for operator `<=>`. Among them is the function `strong_order(const T& a, const T& b)`, specified in section [cmp.alg].

This is the specification in that paper:

```
template<class T>
constexpr strong_ordering strong_order(const T& a, const T& b);
```

1. Effects: Compares two values and produces a result of type `strong_ordering`:

    1.1. If `numeric_limits<T>::is_iec559` is true, returns a result of type `strong_ordering` that is consistent with the totalOrder operation as specified in ISO/IEC/IEEE 60559.

    1.2. Otherwise, returns a `<=>` b if that expression is well-formed and convertible to `strong_ordering`.

1.3. Otherwise, if the expression `a <=> b` is well-formed, then the function shall be defined as deleted.

1.4. Otherwise, if the expressions `a == b` and `a < b` are each well-formed and convertible to `bool`,
returns `strong_ordering::equal` when `a == b` is `true`,
otherwise returns `strong_ordering::less` when `a < b` is `true`,
and otherwise returns `strong_ordering::greater`.

1.5. Otherwise, the function shall be defined as deleted.

Point 1.1 hints at the potential for `strong_order` to be the elusive *default ordering*, required by Stepanov and McJones for *Regular* types, to enable sorting for the purposes of speeding up further processing such as logarithmic searching ([2], page 62, section 4.4). For this purpose, *any* strong ordering would do.

Elements of Programming stresses that many types do not have a *natural order*; even then, a *default order* (a total order that respects at least representational equality) should be provided for all *Regular* types, because the efficiency gains enabled by sorting are enormous. For types that do have a natural total order (possibly only in some of the domain), they specify the *default order* should agree with it wherever defined.

As an example, consider the lexicographic ordering of the gaussian integers. This forms a total order, and its restriction to the integers (gaussian integers of the form $n + 0j$ agrees with the natural order on the integers.

Unfortunately, the hope of finally having a canonical way of naming the default ordering[1] is destroyed by Point 1.3.

Operator `<=>` seems to be explicitly designed[2] to represent the *natural ordering* over the values of `T`. In the case of floating point, iec559 extends this natural order to a total order, thus providing our fabled *default ordering*. However, as per point 1.3, the user is not allowed to specify this extension to `strong_order` themselves, because the function is specified as deleted – it still participates in overload resolution.

3 Code Example

Let me illustrate on a trivial example. Say we have a template struct representing the gaussian integers, which are comparable only if their imaginary part is 0. This constitutes a partial order.

```
namespace user {
template <typename T>
struct gaussian {
  static_assert(std::is_integral_v<T>);
  T real;
  T imag;

  // vector-like
  size_t size() const { return 2; }
  T operator[](size_t i) const {
    assert(i < size());
    return i?imag:real;
  }
};

// our library has this type trait...
template <typename T>
struct is_vector_like : std::false_type {};
template <typename T>
```

---

[1] `std::less` was never really the canonical way of referencing the default ordering, except for pointers.

[2] Because of the various orderings that it supports; they map out the semantic gamut of natural orderings of value types.

```
20   struct is_vector_like<gaussian<T>> : std::true_type {};
21   }
```

We can now define its <=> operator:

```
1    // "natural" ordering on the gaussian integers
2    #include <compare>
3    namespace user {
4    template <typename T>
5    std::partial_ordering operator<=>(gaussian<T> const& x, gassian<T> const& y) {
6      if (x.imag != 0 || y.imag != 0) {
7        return std::partial_ordering::unordered;
8      }
9      return x.real <=> y.real;
10   }
11   }
```

However, since it is extremely useful to still have *some* total order, we shall define the strong_order function template:

```
1    namespace user {
2    template <typename T>
3    std::strong_ordering strong_order(gaussian<T> const& x, gassian<T> const& y) {
4      // compare lexicographically
5      return std::tie(x.real, x.imag) <=> std::tie(y.real, y.imag);
6    }
7    }
```

This works, by virtue of the fact that the template is more specialized than the std::strong_order version.

However, consider overloading on a type trait:

```
1    namespace user {
2    template <typename T>
3    using spaceship_type =
4      std::decay_t<decltype(std::declval<T>() <=> std::declval<T>())>;
5
6    // trait that will disqualify std::strong_order
7    template <typename T>
8    using spaceship_is_not_strong =
9      std::bool_constant<!(std::is_convertible_to<spaceship_type<T>,
10                                                  std::strong_ordering>)>;
11
12
13   // provide a lexicographical order for all vector-like types that define the trait
14   template <typename T,
15             std::enable_if_t<(spaceship_is_not_strong<T>::value &&
16                               is_vector_like<T>::value), void*> = nullptr>
17   std::strong_ordering strong_order(const T & x, const T & y) {
18     using std::strong_order;
19     assert(x.size() == y.size());
20     for (size_t i = 0, s = x.size(); i < s; ++i) {
21       auto const result = strong_order(x[i], y[i]); // ADL
22       if (!std::is_eq(result)) { return result; }
23     }
24     return std::strong_ordering::equal;
25   }
26   }
```

Unfortunatley, this definition is ambiguous with std::strong_order, because std::strong_order is defined as *deleted*, and not as "does not participate in overload resolution".

As an illustration of why having an arbitrary total order available, consider the rest of this example.

Since the standard doesn't yet have an equivalent to std::less and std::equal for std::strong_order, let's roll our own:

```
1   struct strong {
2     struct less {
3       template <typename T>
4       bool operator()(T const& x, T const& y) {
5         using std::strong_order; // use ADL
6         return std::is_lt(strong_order(x, y));
7       }
8     };
9     struct equal {
10      template <typename T>
11      bool operator()(T const& x, T const& y) {
12        using std::strong_order; // use ADL
13        return std::is_eq(strong_order(x, y));
14      }
15    };
16    // also greater, ge_eq and less_eq...
17  };
```

This allows us to unique a vector of any type whatsoever that provides a `strong_order`, even if the type doesn't even provide a comparison operator (or provides a weak or partial one).

```
1   std::vector<gaussian<int>> gaussians = {{1, 0}, {1, 2}, {-1, 2}, {1, 2}};
2   std::sort(gaussians.begin(), gaussians.end(), strong::less{}):
3   gaussians.erase(std::unique(gaussians.begin(), gaussians.end(), strong::equal{}));
```

It allows us to make a `std::set` of these gaussians as well!

```
1   std::set<gaussian<int>, strong::less> gaussian_set = {{1, 0}, {1, 2}, {-1, 2}, {1, 2}}
```

It should be clear that having a canonical way of shipping an arbitrary total order with otherwise unordered, partially ordered or weakly-ordered types is extremely useful for writing efficient generic algorithms.

At last, having a customization point that explicitly says "this is a strong order on this type" is within reach.

4   Proposal

This paper proposes changing point 1.3 to read:

Otherwise, if the expression a <=> b is well-formed,[3] *the function does not participate in overload resolution.*

After the list, add a Note:

If operator <=> provides an order weaker than strong, this function allows the provision of a default strong order for a user-defined type. In that case, `strong_order` *should define a strict, total ordering compatible with the weaker ordering, that is, if a and b are comparable and compare unequal under <=>,* `strong_order` *produces the same result (less or greater).* [4] [5]

5   Fixups

The algorithms section contains a few other algorithms:[6]

---

[3]Note: point 1.2 already takes care of the case where <=> provides a strong (and thus valid default) order.

[4]An ordering weaker than *strong* implies the existence of elements that are unequal, but are not distinguished (deemed either equivalent or unordered) by <=>.

[5]*unequal* does not refer to `operator==`, but to the notion of equality exemplified by iec559 for floating point. It's the answer to the question "what should `unique` do (by default)?".

[6]Not to be confused with the types of their results; those end in -ing: `strong_ordering`, `weak_ordering` etc.

- `weak_order(const T& a, const T& b)`

- `partial_order(const T& a, const T& b)`

- `strong_equal(const T& a, const T& b)`

- `weak_equal(const T& a, const T& b)`

- `partial_equal(const T& a, const T& b)`

Intuitively, one would expect that if `strong_order` is available, then so are `strong_equal`, `weak_order` and `partial_order` (with `weak_equal` and `partial_equal` being consequences of those). The current situation seems to provide for that by pure accident[7], with no reference to this fact.

However, if `strong_order` is the customization point for a default order that *may* be stronger than the order on operator `<=>`, then the above expectation may longer hold for such types.

The fix-up for each of the sections describing the above primitives would be to insert, after point x.1 (which describes the algorithm in terms of `<=>`) the automatic fallback to a call to `strong_order`, if it is resolvable through an unqualified call (thus enabling argument-dependent lookup).

6   Alternative

If the purpose of `strong_order` is not enabling a default-ordering for types, the iec559 exception should be removed from the wording, and a different customization point (perhaps called `total_order`) added for the express purpose of providing an arbitrary total order on the entire domain of a type.

7   Acknowledgments

I would like to thank Thomas Köppe for his valuable comments, review, and most of all some extremely clear and laconic wording; Roger Orr for bringing this to my attention; Sam Finch for *thoroughly* breaking my example and great substantive comments, and pointing out that the current definition actually breaks types that define a partially-ordered set of comparison operators.

References

[1]   Walter E. Brown. "Library Support for the Spaceship (Comparison) Operator". In: *Post-Albuquerque Mailing* (2017). URL: http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/p0768r1.pdf.

[2]   Alexander Stepanov and Paul McJones. *Elements of Programming*. 1st. Addison-Wesley Professional, 2009. ISBN: 032163537X, 9780321635372.

---

[7]the rules are identical, except for the iec559 exception in `strong_order`, while floating-point types possess `operator<` and `operator==`, thus enabling the presence of all those primitives.